

NAME:GANESH KUMAR KORRA

STUDENT ID:700761716

Github Link: [GaneshKumarKorra/ICP4\(github.com\)](https://GaneshKumarKorra/ICP4(github.com))

The screenshot shows a Google Colab notebook titled "Untitled0.ipynb". The code cell contains the following Python code:

```
# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation="relu", kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lrate = 0.01
decay = lrate / epochs
sgd = SGD(learning_rate=lrate, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

Executing (1m 28s) <cell line: 63> > error_handler() > fit() > error_handler() > _call__() > call_function() > _call_flat() > call_preflattened() > call_flat() > call_function() > quick_execute()
```

The status bar at the bottom indicates "73°F Mostly cloudy" and the date/time "7/10/2024 10:51 PM".

The screenshot shows a Google Colab notebook titled "Untitled0.ipynb". The code cell contains the following Python code:

```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from keras.constraints import MaxNorm
from keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

# fix random seed for reproducibility
seed = 7
np.random.seed(seed)

# load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0

# one hot encode outputs
y_train = to_categorical(y_train) # Use to_categorical directly
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation="relu", kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lrate = 0.01
decay = lrate / epochs
sgd = SGD(learning_rate=lrate, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

Executing (1m 28s) <cell line: 63> > error_handler() > fit() > error_handler() > _call__() > call_function() > _call_flat() > call_preflattened() > call_flat() > call_function() > quick_execute()
```

The status bar at the bottom indicates "Finance headline China's GDP rec..." and the date/time "7/10/2024 10:59 PM".

Untitled0.ipynb - Colab

colab.research.google.com/drive/1zo1Ntnb1z3fhS8KNLAp0vU3mnfr8gT?usp=chrome_ntp

Untitled0.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
# compile model
epochs = 25
lrate = 0.01
decay = lrate / epochs
sgd = SGD(learning_rate=lrate, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
predicted_classes = np.argmax(predictions, axis=1)
actual_classes = np.argmax(y_test[:4], axis=1)

for i in range(4):
    print(f"Image {i+1}: Predicted: {predicted_classes[i]}, Actual: {actual_classes[i]}")

# visualize the first 4 test images
for i in range(4):
    plt.imshow(X_test[i])
    plt.title(f"Predicted: {predicted_classes[i]}, Actual: {actual_classes[i]}")
    plt.show()

# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# Plot loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Plot accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

Executing (1m 57s) <cell line: 63> <error_handler()> fit() <error_handler()> _call_() <call_function()> _call_flat() <call_preflattened()> call_flat() <call_function()> quick_execute()

73°F Mostly cloudy

ENG INTL 10:51 PM 7/10/2024

Untitled0.ipynb - Colab

colab.research.google.com/drive/1zo1Ntnb1z3fhS8KNLAp0vU3mnfr8gT?usp=chrome_ntp

Untitled0.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# plot loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Plot accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

... Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> 170498071/170498071 [=====] - 6s 0us/step Model: "sequential"

Layer (Type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896

denorm (Denorm) (None, 32, 32, 32)

Executing (2m 39s) <cell line: 63> <error_handler()> fit() <error_handler()> _call_() <call_function()> _call_flat() <call_preflattened()> call_flat() <call_function()> quick_execute()

Rain coming 10:52 PM

ENG INTL 10:52 PM 7/10/2024

output

A screenshot of a Google Colab notebook titled "Untitled0.ipynb". The code cell contains the following Python code:

```
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

Below the code is a table showing the memory usage of various layers in the model:

Layer	Type	Shape	Size (bytes)
conv2d_1	conv2d	(None, 32, 32, 32)	9248
max_pooling2d	MaxPooling2D	(None, 16, 16, 32)	0
conv2d_2	conv2d	(None, 16, 16, 64)	18496
dropout_1	Dropout	(None, 16, 16, 64)	0
conv2d_3	conv2d	(None, 16, 16, 64)	36928
max_pooling2d_1	MaxPooling2D	(None, 8, 8, 64)	0
conv2d_4	conv2d	(None, 8, 8, 128)	73856
dropout_2	Dropout	(None, 8, 8, 128)	0
conv2d_5	conv2d	(None, 8, 8, 128)	147584
max_pooling2d_2	MaxPooling2D	(None, 4, 4, 128)	0
flatten	Flatten	(None, 2048)	0
dropout_3	Dropout	(None, 2048)	0
dense	Dense	(None, 1024)	2098176

The status bar at the bottom indicates "Executing (6m 40s) <cell line: 63> > error_handler()".

A screenshot of a Google Colab notebook titled "Untitled0.ipynb". The code cell contains the following Python code:

```
dropout_4 (Dropout) (None, 1024) 0
dense_1 (Dense) (None, 512) 524800
dropout_5 (Dropout) (None, 512) 0
dense_2 (Dense) (None, 10) 5130
```

Below the code is a summary table:

Total params:	2915114 (11.12 MB)
Trainable params:	2915114 (11.12 MB)
Non-trainable params:	0 (0.00 Byte)

Following the summary table is a detailed training log output:

```
None
Epoch 1/25
1563/1563 [=====] - 21s 9ms/step - loss: 1.8506 - accuracy: 0.3148 - val_loss: 1.4900 - val_accuracy: 0.4600
Epoch 2/25
1563/1563 [=====] - 13s 8ms/step - loss: 1.4098 - accuracy: 0.4869 - val_loss: 1.2368 - val_accuracy: 0.5581
Epoch 3/25
1563/1563 [=====] - 13s 8ms/step - loss: 1.1750 - accuracy: 0.5772 - val_loss: 1.0572 - val_accuracy: 0.6181
Epoch 4/25
1563/1563 [=====] - 13s 8ms/step - loss: 1.0207 - accuracy: 0.6360 - val_loss: 0.9175 - val_accuracy: 0.6807
Epoch 5/25
1563/1563 [=====] - 14s 9ms/step - loss: 0.9108 - accuracy: 0.6797 - val_loss: 0.8522 - val_accuracy: 0.7007
Epoch 6/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.8339 - accuracy: 0.7084 - val_loss: 0.8547 - val_accuracy: 0.6997
Epoch 7/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.7716 - accuracy: 0.7270 - val_loss: 0.7809 - val_accuracy: 0.7325
Epoch 8/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.7193 - accuracy: 0.7480 - val_loss: 0.7539 - val_accuracy: 0.7386
Epoch 9/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.6825 - accuracy: 0.7615 - val_loss: 0.7355 - val_accuracy: 0.7467
Epoch 10/25
1563/1563 [=====] - 13s 8ms/step - loss: 0.6441 - accuracy: 0.7743 - val_loss: 0.6887 - val_accuracy: 0.7643
```

The status bar at the bottom indicates "6m 52s completed at 10:56PM".

Untitled0.ipynb - Colab

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
Epoch 11/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.6213 - accuracy: 0.7842 - val_loss: 0.7006 - val_accuracy: 0.7589  
Epoch 12/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5988 - accuracy: 0.7918 - val_loss: 0.6669 - val_accuracy: 0.7723  
Epoch 13/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5846 - accuracy: 0.7958 - val_loss: 0.6776 - val_accuracy: 0.7721  
Epoch 14/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5637 - accuracy: 0.8046 - val_loss: 0.7321 - val_accuracy: 0.7631  
Epoch 15/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5563 - accuracy: 0.8059 - val_loss: 0.7261 - val_accuracy: 0.7539  
Epoch 16/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5550 - accuracy: 0.8086 - val_loss: 0.6751 - val_accuracy: 0.7790  
Epoch 17/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5330 - accuracy: 0.8160 - val_loss: 0.6932 - val_accuracy: 0.7678  
Epoch 18/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5228 - accuracy: 0.8175 - val_loss: 0.7228 - val_accuracy: 0.7562  
Epoch 19/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5299 - accuracy: 0.8174 - val_loss: 0.6796 - val_accuracy: 0.7719  
Epoch 20/25  
1563/1563 [=====] - 13s 9ms/step - loss: 0.5215 - accuracy: 0.8215 - val_loss: 0.7126 - val_accuracy: 0.7599  
Epoch 21/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5290 - accuracy: 0.8172 - val_loss: 0.7075 - val_accuracy: 0.7573  
Epoch 22/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5256 - accuracy: 0.8201 - val_loss: 0.7205 - val_accuracy: 0.7592  
Epoch 23/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5343 - accuracy: 0.8174 - val_loss: 0.7392 - val_accuracy: 0.7589  
Epoch 24/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5409 - accuracy: 0.8161 - val_loss: 0.7242 - val_accuracy: 0.7591  
Epoch 25/25  
1563/1563 [=====] - 13s 8ms/step - loss: 0.5453 - accuracy: 0.8148 - val_loss: 0.6899 - val_accuracy: 0.7755  
Accuracy: 77.55%  
1/1 [=====] - 0s 457ms/step  
Image 1: Predicted: 3, Actual: 3  
Image 2: Predicted: 8, Actual: 8  
Image 3: Predicted: 8, Actual: 8
```

6m 52s completed at 10:56 PM

from W 119th St Construction

ENG INTL 10:58 PM 7/10/2024

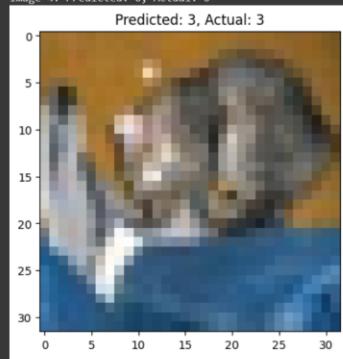
Untitled0.ipynb - Colab

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
Image 3: Predicted: 8, Actual: 8  
Image 4: Predicted: 0, Actual: 0
```

Predicted: 3, Actual: 3



Predicted: 8, Actual: 8



6m 52s completed at 10:56 PM

from Overland P... Clearing soon

ENG INTL 10:59 PM 7/10/2024

